

How to build & run your first deep learning network in TensorFlow

Jordi Torres & Maurici Yagües

Barcelona, July 2016



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

DAY 2

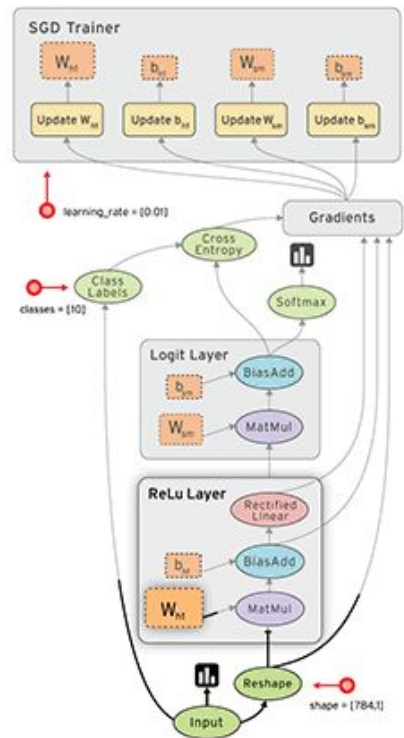
Basic data structures in TensorFlow

Case study: Clustering in TensorFlow

Parallelism and Distribution in TensorFlow

Core TensorFlow concepts (reminder)

- **Graph:** A TensorFlow computation, represented as a dataflow graph
- **Operation:** a graph node that performs computation on tensors
- **Tensor:** a handle to one of the outputs of an operation



Tensor: Core TensorFlow data structure (reminder)

- **Constants**
- **Placeholders**: must be fed with data on execution
- **Variables**: a modifiable tensor that lives in TensorFlow's graph of interacting operations
- **Session**: encapsulates the environment in which operation objects are executed, and Tensor objects are evaluated

Basic data type: Tensor

- Tensor can be considered a dynamically-sized multidimensional data arrays
- Main types and their equivalent in Python:

Type in TensorFlow	Type in Python	Description
DT_FLOAT	tf.float32	Floating point of 32 bits
DT_INT16	tf.int16	Integer of 16 bits
DT_INT32	tf.int32	Integer of 32 bits
DT_INT64	tf.int64	Integer of 64 bits
DT_STRING	tf.string	String
DT_BOOL	tf.bool	Boolean

Basic data type: Tensor

- Each tensor has a rank, which is the number of its dimensions.
- For example, the following tensor (defined as a list in Python) has rank 2:

```
t = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Basic data type: Tensor

- TensorFlow (documentation) uses three types of naming conventions

Shape	Rank	Dimension Number
[]	0	0-D
[D0]	1	1-D
[D0, D1]	2	2-D
[D0, D1, D2]	3	3-D
...
[D0, D1, ... Dn]	n	n-D

Transformations

- Tensors can be manipulated with a series of transformations that supply the TensorFlow package

Operation	Description
tf.shape	To find a shape of a <i>tensor</i>
tf.size	To find the size of a <i>tensor</i>
tf.rank	To find a rank of a <i>tensor</i>
tf.reshape	To change the shape of a <i>tensor</i> keeping the same elements contained
tf.squeeze	To delete in a <i>tensor</i> dimensions of size 1
tf.expand_dims	To insert a dimension to a <i>tensor</i>

Transformations (cont.)

Operation	Description
tf.slice	To remove a portions of a <i>tensor</i>
tf.split	To divide a <i>tensor</i> into several tensors along one dim
tf.tile	To create a new <i>tensor</i> replicating it multiple times
tf.concat	To concatenate <i>tensors</i> in one dimension
tf.reverse	To reverse a specific dimension of a <i>tensor</i>
tf.transpose	To transpose dimensions in a <i>tensor</i>
tf.gather	To collect portions according to an index

Tensor transformations




- For example, suppose that we want to extend an array of 2x2000 (a 2D tensor) to a cube (3D tensor).
- We can use the **tf.expand_dims** function, which allows us to insert a dimension to a tensor (the dimensions start at zero):

```
vectors = tf.constant(vectors_set)
extended_vectors = tf.expand_dims(vectors, 0)
```

Tensor transformations

- For example, suppose that we want to extend an array of 2x2000 (a 2D tensor) to a cube (3D tensor).
- We can use the `tf.expand_dims` function, which allows us to insert a dimension to a tensor (the dimensions start at zero):

```
vectors = tf.constant(vectors_set)
extended_vectors = tf.expand_dims(vectors, 0)
```




Tensor shape

- If we obtain the shape of this tensor with the `.get_shape()` operation:

```
print expanded_vectors.get_shape()
```

- we can see that there is no associated size:

```
TensorShape([Dimension(1), Dimension(2000), Dimension(2)])
```



expanded_vectors

Data storage in TensorFlow

There are three main ways of obtaining data on a TensorFlow program:

1. From data files.
2. Data preloaded as constants or variables.
3. Those provided by Python code.

Data storage in TensorFlow

1. From data files: example `input_data.py`

```
SOURCE_URL = 'http://yann.lecun.com/exdb/mnist/'  
TRAIN_IMAGES = 'train-images-idx3-ubyte.gz'
```

```
local_file = maybe_download(TRAIN_IMAGES, train_dir)  
train_images = extract_images(local_file)
```

```
def maybe_download(filename, work_directory):  
    """Download the data from Yann's website, unless it's already here."""  
    if not os.path.exists(work_directory):  
        os.mkdir(work_directory)  
    filepath = os.path.join(work_directory, filename)  
    if not os.path.exists(filepath):  
        filepath, _ = urllib.request.urlretrieve(SOURCE_URL + filename, filepath)  
        statinfo = os.stat(filepath)  
        print('Successfully downloaded', filename, statinfo.st_size, 'bytes.')  
    return filepath
```

```
def extract_images(filename):  
    """Extract the images into a 4D uint8 numpy array [index, y, x, depth]."""  
    print('Extracting', filename)  
    with gzip.open(filename) as bytestream:  
        magic = _read32(bytestream)  
        if magic != 2051:  
            raise ValueError(  
                'Invalid magic number %d in MNIST image file: %s' %  
                (magic, filename))  
        num_images = _read32(bytestream)  
        rows = _read32(bytestream)  
        cols = _read32(bytestream)  
        buf = bytestream.read(rows * cols * num_images)  
        data = numpy.frombuffer(buf, dtype=numpy.uint8)  
        data = data.reshape(num_images, rows, cols, 1)  
    return data
```

Data storage in TensorFlow

2. Data preloaded as a:

- Constant using `tf.constant(...)`
- Variable using `tf.Variable(...)`

TensorFlow package offers different operations that can be used to generate constants and variables:

Constants generation

Operation	Description
<code>tf.zeros_like</code>	Creates a tensor with all elements initialized to 0
<code>tf.ones_like</code>	Creates a tensor with all elements initialized to 1
<code>tf.fill</code>	Creates a tensor with all elements initialized to a scalar value given as argument
<code>tf.constant</code>	Creates a tensor of constants with the elements listed as an arguments

Variable Tensor random generation

Operation	Description
<code>tf.random_normal</code>	Random values with a normal distribution
<code>tf.truncated_normal</code>	Random values with a normal distribution but eliminating those values whose magnitude is more than 2 times the standard deviation
<code>tf.random_uniform</code>	Random values with a uniform distribution
<code>tf.random_shuffle</code>	Randomly mixed tensor elements in the first dimension
<code>tf.set_random_seed</code>	Sets the random seed

Data storage in TensorFlow

3. Provided by Python code:

The call is **`tf.placeholder()`**, which includes arguments with the type of the elements and the shape of the tensor, and optionally a name.

```
import tensorflow as tf

a = tf.placeholder("float")
b = tf.placeholder("float")

y = tf.mul(a, b)

sess = tf.Session()
print sess.run(y, feed_dict={a: 3, b: 3})
```

*With calls in `Session.run()` or `Tensor.eval()` this tensor is populated with the data specified in the **`feed_dict`** parameter*

Case study: K-means algorithm

- K-means: Unsupervised algorithm which solves the clustering problem
 - Its procedure follows a simple and easy way to classify a given data set through a certain number of clusters (**assume k clusters**).
 - Data points inside a cluster are homogeneous and heterogeneous to peer groups, that means that all the elements in a subset are more similar to each other than with the rest.
 - The result of the algorithm is a set of K dots, called **centroids**, which are the focus of the different groups obtained, and the tag that represents the set of points that are assigned to only one of the K clusters. All the points within a cluster are closer in distance to the centroid than any of the other centroids.


Case study: K-means algorithm

- K-means:
 - Computationally expensive problem (NP-hard problem)
 - Using algorithms that converge rapidly in a local optimum by heuristics
 - The most commonly used algorithm uses an iterative refinement technique.
- Steps:
 1. Initial step: **determines an initial set of K centroids.**
 2. Allocation step: **assigns each observation to the nearest group.**
 3. Update step: **calculates the new centroids for each new group.**
 4. Steps 2 and 3 are repeated until convergence has been reached.

k-means algorithm: visual example

1. Initial step: **determines an initial set of K centroids.**


k initial “means” (in this case $k = 3$) are randomly generated within the data domain (shown in color)



k-means algorithm: visual example

2. Allocation step: **assigns each observation to the nearest group.**


k clusters are created by associating every observation with the nearest mean.
The partitions here represent Voronoi diagram generated by the means



k-means algorithm: visual example


3. Update step: **calculates the new centroids for each new group.**

The centroid of each of the k clusters becomes the new mean. k clusters are created by associating every observation with the nearest mean



k-means algorithm: visual example

4. Steps 2 and 3 are repeated until convergence has been reached



- As it is a heuristic algorithm, there is no guarantee that it will converge to the global optimum, and the result may depend on the initial clusters
- As the algorithm is usually very fast, it is common to run it multiple times with different starting conditions


K-means: one testbed example

- *Generate 2000 points in a 2D space in a random manner (following 2 normal distributions)*

```
num_puntos = 2000
conjunto_puntos = []

for i in xrange(num_puntos):
    if np.random.random() > 0.5:
        conjunto_puntos.append([np.random.normal(0.0, 0.9),
                                np.random.normal(0.0, 0.9)])
    else:
        conjunto_puntos.append([np.random.normal(3.0, 0.5),
                                np.random.normal(1.0, 0.5)])
```

K-means: one testbed example



```
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

df = pd.DataFrame({"x": [v[0] for v in vectors_set],
                  "y": [v[1] for v in vectors_set]})
sns.lmplot("x", "y", data=df, fit_reg=False, size=6)
plt.show()
```

K-means: one testbed example (step by step)

1. Initial step: determines an initial set of K centroids

- The first thing to do is move all our data to tensors:

```
vectors = tf.constant(vectors_set)
```

- Randomly choose K observations from the input data as centroids

```
k = 4
```

```
centroids = tf.Variable(tf.slice(tf.random_shuffle(vectors), [0,0],  
[k,-1]))
```

```
(*)centroids.get_shape() → TensorShape([Dimension(4), Dimension(2)])
```

K-means: one testbed example (step by step)

2. Allocation step: assigns each observation to the nearest group

- Calculate, for each point, its closest centroid by the *Squared Euclidean Distance*:



$$d^2(\text{vector}, \text{centroid}) = (\text{vector}_x - \text{centroid}_x)^2 + (\text{vector}_y - \text{centroid}_y)^2$$

- TensorFlow code ?

K-means: one testbed example (step by step)

(2. Allocation step: assigns each observation to the nearest group)

- `tf.sub(vectors, centroids)` is the main function used.
 - note that, although the two subtract tensors have both 2 dimensions, they have different sizes in one dimension (2000 vs 4 in dimension D0), which, in fact, also represent different things.



K-means: one testbed example (step by step)


- To fix this problem we could use some of the functions discussed before, for instance **tf.expand_dims** in order to insert a dimension in both tensors.
- The aim is to extend both tensors from 2 dimensions to 3 dimensions to make the sizes match in order to perform a subtraction:

```
expanded_vectors = tf.expand_dims(vectors, 0)
expanded_centroides = tf.expand_dims(centroides, 1)
```

K-means: one testbed example (step by step)

`tf.expand_dims`

- inserts one dimension in each tensor; in the first dimension (D0) of vectors tensor,
- and in the second dimension (D1) of centroids tensor.



K-means: one testbed example (step by step)

Note: there are dimensions that have not been able to determinate the sizes of those dimensions:

```
print expanded_vectors.get_shape()  
print expanded_centroides.get_shape()
```

```
TensorShape([Dimension(1), Dimension(2000), Dimension(2)])  
TensorShape([Dimension(4), Dimension(1), Dimension(2)])
```

() With 1 it is indicating a no assigned size.*

TensorFlow allows **broadcasting**, and therefore the **tf.sub** function is able to discover for itself how to do the subtraction of elements between the two tensors.

K-means: one testbed example (step by step)

$$d^2(\text{vector}, \text{centroide}) = (\text{vector}_x - \text{centroide}_x)^2 + (\text{vector}_y - \text{centroide}_y)^2$$

- TensorFlow code ?

```
expanded_vectors = tf.expand_dims(vectors, 0)
expanded_centroides = tf.expand_dims(centroides, 1)

diff=tf.sub(expanded_vectors, expanded_centroides)
sqr= tf.square(diff)
```

(*) tensor diff shape is **TensorShape([Dimension(4), Dimension(2000), Dimension(2)])** where D0 indicates the centroid , D1 the subtraction value and D2 each x,y point. sqr have the same shape.

K-means: one testbed example (step by step)

2. Allocation step: assigns each observation to the nearest group.


- Code that calculates, for each point, its closest centroid:

```
distances = tf.reduce_sum(sqr, 2)  
assignments = tf.argmin(distances, 0)
```

K-means: one testbed example (step by step)

```
distances = tf.reduce_sum(sqr, 2)
```

- The **distance** tensor has already reduced one dimension, the one indicated as a parameter in `tf.reduce_sum` function `TensorShape([Dimension(4), Dimension(2000)])`




K-means: one testbed example (step by step)

```
assignments = tf.argmin(distances, 0)
```

- The assignment is achieved with `tf.argmin`, which returns the index with the minimum value of the tensor dimension (in our case `D0`, which remember that was the centroid). `TensorShape([Dimension(2000)])`

K-means: one testbed example (step by step)

```
distances = tf.reduce_sum(sqr, 2)
assignments = tf.argmin(distances, 0)
```



Tensor operations

Operation	Description
<code>tf.reduce_sum</code>	Computes the sum of elements along one dimensions of a tensor
<code>tf.reduce_prod</code>	Computes the product of elements along one dimensions of a tensor
<code>tf.reduce_min</code>	Computes the minimum of elements along one dimensions of a tensor
<code>tf.reduce_max</code>	Computes the maximum of elements along one dimensions of a tensor
<code>tf.reduce_mean</code>	Computes the mean of elements along one dimensions of a tensor

Tensor operations

Operation	Description
<code>tf.argmin</code>	Returns the index of the element with the minimum value along tensor dimension
<code>tf.argmax</code>	Returns the index of the element with the maximum value of the tensor dimension

K-means: one testbed example (step by step)

```
diff=tf.sub(expanded_vectors, expanded_centroides)
sqr= tf.square(diff)
distances = tf.reduce_sum(sqr, 2)
assignments = tf.argmin(distances, 0)
```

- Alternative TensorFlow code used in my book (*)

```
assignments = tf.argmin(tf.reduce_sum(tf.square(tf.sub
(expanded_vectors,
      expanded_centroides))), 2), 0)
```

(*) Github, (2016) Shawn Simister. [Online]. Available at: <https://gist.github.com/narphorium/d06b7ed234287e319f18>

K-means: one testbed example (step by step)

3. Update step: calculates the new centroids for each new group

- We create a tensor that contains the result of the concatenation of the k tensors that correspond to the mean value of every point that belongs to each k cluster (*):

```
means = tf.concat(0, [tf.reduce_mean(tf.gather(vectors, tf.reshape(
    tf.where( tf.equal(assignments, c)), [1,-1])), reduction_indices=[1])
    for c in xrange(k)])
```

(*) Github, (2016) Shawn Simister. [Online]. Available at: <https://gist.github.com/narphorium/d06b7ed234287e319f18>

(brief explanation)

- With `tf.equal` we can obtain a **boolean tensor** (*Dimension(2000)*) that indicates (with *true* value) the positions where the *assignment tensor* match with the *K cluster*, which, at the time, we are calculating the average value of the points.
- With `tf.where` is constructed a tensor (*Dimension(1) x Dimension(2000)*) with the position where the values *true* are on the *boolean tensor* received as a parameter. (i.e. a list of the position of these)
- With `tf.reshape` is constructed a tensor (*Dimension(2000) x Dimension(1)*) with the index of the points inside *vectors* tensor that belongs to this *c cluster*
- With `tf.gather` is constructed a *tensor* (*Dimension(1) x Dimension(2000)*) which gathers the coordinates of the points that form the *c cluster*
- With `tf.reduce_mean` it is constructed a tensor (*Dimension(1) x Dimension(2)*) that contains the average value of all points that belongs to the cluster *c*

K-means: one testbed example

4. Graph Execution

```
update_centroides = tf.assign(centroides, means)

init_op = tf.initialize_all_variables()

sess = tf.Session()
sess.run(init_op)

for step in xrange(100):
    _, centroid_values, assignment_values = sess.run([update_centroides,
                                                       centroides, assignments])
```

K-means: one testbed example

```
vectors = tf.constant(conjunto_puntos)
k = 4
centroides = tf.Variable(tf.slice(tf.random_shuffle(vectors), [0,0], [k,-1]))

expanded_vectors = tf.expand_dims(vectors, 0)
expanded_centroides = tf.expand_dims(centroides, 1)

assignments = tf.argmin(tf.reduce_sum(tf.square(tf.sub(expanded_vectors,
                                                         expanded_centroides)), 2), 0)

means = tf.concat(0, [tf.reduce_mean(tf.gather(vectors, tf.reshape(tf.where( tf.equal
(assignments, c)),[1,-1])), reduction_indices=[1]) for c in xrange(k)])

update_centroides = tf.assign(centroides, means)

init_op = tf.initialize_all_variables()

sess = tf.Session()
sess.run(init_op)

for step in xrange(100):
    _, centroid_values, assignment_values = sess.run([update_centroides,
                                                         centroides, assignments])
```

K-means: one testbed example


- We can use a simple print to know the centroids:

```
print centroid_values
```

- The output is as it follows:

```
[[ 2.99835277e+00  9.89548564e-01]
 [-8.30736756e-01  4.07433510e-01]
 [ 7.49640584e-01  4.99431938e-01]
 [ 1.83571398e-03 -9.78474259e-01]]
```

Case study: k-means



```
data = {"x": [], "y": [], "cluster": []}

for i in xrange(len(assignment_values)):
    data["x"].append(conjunto_puntos[i][0])
    data["y"].append(conjunto_puntos[i][1])
    data["cluster"].append(assignment_values[i])

df = pd.DataFrame(data)
sns.lmplot("x", "y", data=df,
           fit_reg=False, size=6,
           hue="cluster", legend=False)

plt.show()
```

Some words about Parallelism in TensorFlow

- The TensorFlow package, appearing in November 2015, was ready to run on servers with available GPUs and executing the training operation simultaneously in them
 - Requires the CudaToolkit and CUDNN packages
- In February 2016, an update added the capability to distribute and parallelize the processing

GPUs

- The way to reference those devices in TensorFlow is the following one:
 - `"/cpu:0"`: To reference the server's CPU.
 - `"/gpu:0"`: The server's GPU, if only one is available.
 - `"/gpu:1"`: The second server's GPU, and so on
- To know (in the output) in which devices our operations and tensors are assigned we need to create a session with the option `log_device_placement` as `True`.

```
sess = tf.Session(config=tf.ConfigProto(log_device_placement=True))
```

- If we want a specific operation to be executed in a specific device:
`tf.device('/gpu:2')`

Example of GPU use:

```
import tensorflow as tf

with tf.device('/gpu:2'):
    a = tf.constant([1.0, 2.0, 3.0, 4.0, 5.0, 6.0], shape=[2, 3], name='a')
    b = tf.constant([1.0, 2.0, 3.0, 4.0, 5.0, 6.0], shape=[3, 2], name='b')
    c = tf.matmul(a, b)

sess = tf.Session(config=tf.ConfigProto(log_device_placement=True))
print sess.run(c)
```

Example with two GPUs :

```
import tensorflow as tf

c = []
for d in ['/gpu:2', '/gpu:3']:
    with tf.device(d):
        a = tf.constant([1.0, 2.0, 3.0, 4.0, 5.0, 6.0], shape=[2, 3])
        b = tf.constant([1.0, 2.0, 3.0, 4.0, 5.0, 6.0], shape=[3, 2])
        c.append(tf.matmul(a, b))
with tf.device('/cpu:0'):
    sum = tf.add_n(c)

sess = tf.Session(config=tf.ConfigProto(log_device_placement=True))
print sess.run(sum)
```

Output using log_device_placement=True

```
. . .
Device mapping:
/job:localhost/replica:0/task:0/gpu:0 -> device: 0, name: Tesla K40c
/job:localhost/replica:0/task:0/gpu:1 -> device: 1, name: Tesla K40c
/job:localhost/replica:0/task:0/gpu:2 -> device: 2, name: Tesla K40c
/job:localhost/replica:0/task:0/gpu:3 -> device: 3, name: Tesla K40c
. . .

Const_3: /job:localhost/replica:0/task:0/gpu:3
I tensorflow/core/common_runtime/simple_placer.cc:289] Const_3: /job:localhost/replica:0/task:0/gpu:3
Const_2: /job:localhost/replica:0/task:0/gpu:3
I tensorflow/core/common_runtime/simple_placer.cc:289] Const_2: /job:localhost/replica:0/task:0/gpu:3
MatMul_1: /job:localhost/replica:0/task:0/gpu:3
I tensorflow/core/common_runtime/simple_placer.cc:289] MatMul_1: /job:localhost/replica:0/task:0/gpu:3
Const_1: /job:localhost/replica:0/task:0/gpu:2
I tensorflow/core/common_runtime/simple_placer.cc:289] Const_1: /job:localhost/replica:0/task:0/gpu:2
Const: /job:localhost/replica:0/task:0/gpu:2
I tensorflow/core/common_runtime/simple_placer.cc:289] Const: /job:localhost/replica:0/task:0/gpu:2
MatMul: /job:localhost/replica:0/task:0/gpu:2
I tensorflow/core/common_runtime/simple_placer.cc:289] MatMul: /job:localhost/replica:0/task:0/gpu:2
AddN: /job:localhost/replica:0/task:0/cpu:0
I tensorflow/core/common_runtime/simple_placer.cc:289] AddN: /job:localhost/replica:0/task:0/cpu:0
[[44.56.]
 [98.128.]]
. . .
```

Distributed TensorFlow: gRPC

More information: https://www.tensorflow.org/versions/r0.8/how_tos/distributed/index.html

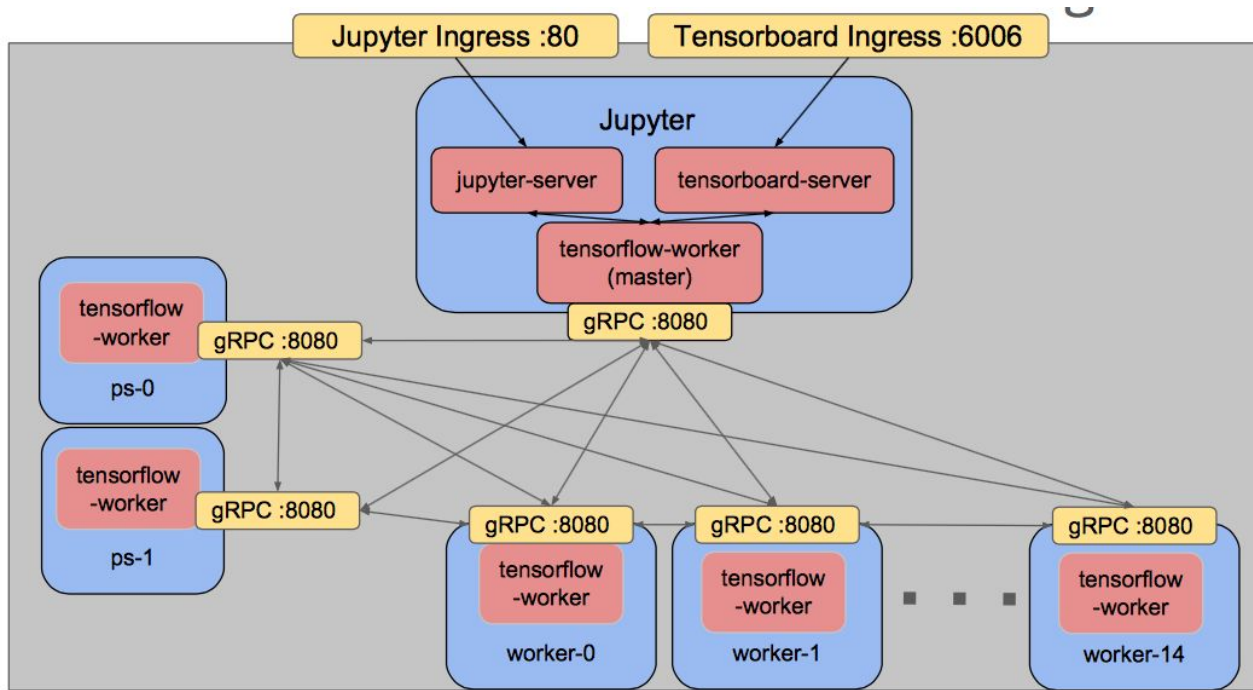


Image source: <https://storage.googleapis.com/amy-jo/talks/tf-workshop.pdf>

Class hands-on (or Homework)

- Reproduce the results (you can download the code from github) presented in this course in your laptop and show them to the instructors.

Next session:

- We will build a single layer neural network, step by step, with TensorFlow and use TensorBoard, a graph visualization tool