

How to build & run your first deep learning network in TensorFlow

Jordi Torres & Maurici Yagües

Barcelona, July 2016



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación


DAY 3

Single Layer Neural Network in TensorFlow

TensorBoard

Single Layer Neural Network in TensorFlow

- MNIST Data-set : “Hello World”
 - set of black and white images containing hand-written digits, containing more than 60.000 examples for training a model, and 10.000 for testing it.
 - The images have been normalized into 28x28 pixel images, preserving the aspect ratio. After that, the images are centered in 28x28 pixel frames
 - The images are labeled with the digit they represent.
 - The images are like:



Single Layer Neural Network in TensorFlow

- MNIST Data-set : “Hello World”
 - To download easily the data, you can use the script `input_data.py`
 - From your application you only need to import and use in the following way:

```
import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

Single Layer Neural Network in TensorFlow

- MNIST Data-set : “Hello World”
 - After executing these two instructions you will have
 - the full training data-set in *mnist.train* and
 - the test data-set in *mnist.test*.
 - each element is composed by an image, referenced as “xs”, and its corresponding label “ys”

Single Layer Neural Network in TensorFlow

- An easy example to start: Softmax

The *softmax* function has two main steps:

1. Compute the “evidences” for an image belonging to a certain label

$$evidence_i = \sum_j W_{ij} x_j + b_i$$



2. Convert the evidences into probabilities for each possible label

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)} = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

Single Layer Neural Network in TensorFlow

1. Compute the “evidences” for an image belonging to a certain label are

- a usual approximation is to compute the weighted sum of pixel intensities.
- That weight is negative when a pixel with high intensity happens to not to be in a given class,
- and positive if the pixel is frequent in that class.




Single Layer Neural Network in TensorFlow

1. Compute the “evidences” for an image belonging to a certain label are

- In a more formal way, we can say that the evidence for a class i given an input x is expressed as:

$$\text{evidence}_i = \sum_j W_{ij} X_j$$




Single Layer Neural Network in TensorFlow

1. Compute the “evidences” for an image belonging to a certain label are

- The model also include an extra parameter representing the bias (adding some base uncertainty)

$$evidence_i = \sum_j W_{ij} X_j + b_i$$



Single Layer Neural Network in TensorFlow

2. Convert the evidences into probabilities for each possible label.


$$y = \text{softmax}(\text{evidence})$$

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)} = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

The interesting fact of such function is that a good prediction will have one output with a value near 1, while all the other outputs will be near zero; and in a weak prediction, some labels may show similar support.

Single Layer Neural Network in TensorFlow


Data structure:



Two variable are created using the tf.Variable function

```
W = tf.Variable(tf.zeros([784,10]))  
b = tf.Variable(tf.zeros([10]))
```

Single Layer Neural Network in TensorFlow




We create a tensor of two dimensions to keep the information of the x points (used to store the MNIST images as a vector of 784 floating point values (*)).

```
x = tf.placeholder("float", [None, 784])
```

(*) None indicates that the dimension can be any size

Single Layer Neural Network in TensorFlow




`tf.matmul(x, W) + b`

Single Layer Neural Network in TensorFlow

- implementing the previously described softmax function:

```
y = tf.nn.softmax(tf.matmul(x,W) + b)
```



Single Layer Neural Network in TensorFlow

- Once specified the model implementation, we can specify the necessary code to obtain the weights for W and bias b using an iterative training algorithm
 - For each iteration, the training algorithm gets the training data, applies the neural network and compares the obtained result with the expected one.

- As a cost function : *cross entropy error*

$$-\sum_i y_i' \log(y_i)$$

Single Layer Neural Network in TensorFlow

- To implement the cross-entropy measurement we need:

```
y_ = tf.placeholder("float", [None,10])  
cross_entropy = -tf.reduce_sum(y_*tf.log(y))
```

- As a iterative minimization process: backpropagation with gradient descent method using the cross-entropy cost function.

```
train_step = tf.train.GradientDescentOptimizer(0.01).minimize  
(cross_entropy)
```

Single Layer Neural Network in TensorFlow

- Start the computation by instantiating `tf.Session()`:

```
sess = tf.Session()

sess.run(tf.initialize_all_variables())

for i in range(1000):

    batch_xs, batch_ys = mnist.train.next_batch(100)

    sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
```

For each iteration, a bundle of 100 inputs of data, randomly sampled from the training data-set, are picked.

The returning parameter for *train_step*, when executed, will apply the gradient descent to the involved parameters. So training the model can be achieved by repeating the *train_step* execution.

Single Layer Neural Network in TensorFlow

- A **model must be evaluated** after training to see how much “good” is!
 - For example, we can compute the percentage of hits and misses in our prediction, seeing which examples were correctly predicted

```
correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
```

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```

```
print sess.run(accuracy, feed_dict={x: mnist.test.images, y_:  
mnist.test.labels})
```

Single Layer Neural Network in TensorFlow

```
correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
```

the *tf.argmax(y, 1)* function returns the index of the highest value of a tensor according a given axis. In effect, *tf.argmax(y, 1)* is the label in our model with higher probability for each input, while *tf.argmax(y_, 1)* is the correct label.

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```

Using *tf.equal* method we can compare if our prediction coincides with the correct label.

For example, *[True, False, True, True]* will turn into *[1,0,1,1]* and the average will be 0.75 representing the percentage of *accuracy*. Now we can ask for the accuracy of our test data-set using the *mnist.test* as the *feed_dict* argument:

Single Layer Neural Network in TensorFlow

Using *tf.equal* method we can compare if our prediction coincides with the correct label:

```
correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
```

This instruction returns a list of Booleans. To determine which fractions of predictions are correct, we can cast the values to numeric variables (floating point) and do the following operation:

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```

For example, *[True, False, True, True]* will turn into *[1,0,1,1]* and the average will be 0.75 representing the percentage of *accuracy*. Now we can ask for the accuracy of our test data-set using the *mnist.test* as the *feed_dict* argument:


```
print sess.run(accuracy, feed_dict={x: mnist.test.images, y_: mnist.test.labels})
```

Class hands-on (or Homework)

- Download the code [SingleLayerNeuralNetwork.py](#) from github and indicates the obtained value.

TensorBoard

- Built-in graph visualization tool
- Can help you better understand the model
- Show and plot quantitative metrics about your graph
- Requires some extra work on the original code
- Some links for reference:
 - [TensorBoard: Visualizing Learning](#)
 - [TensorBoard: Graph Visualization](#)
 - [Running demo \(code\)](#)
 - [TensorBoard README](#)



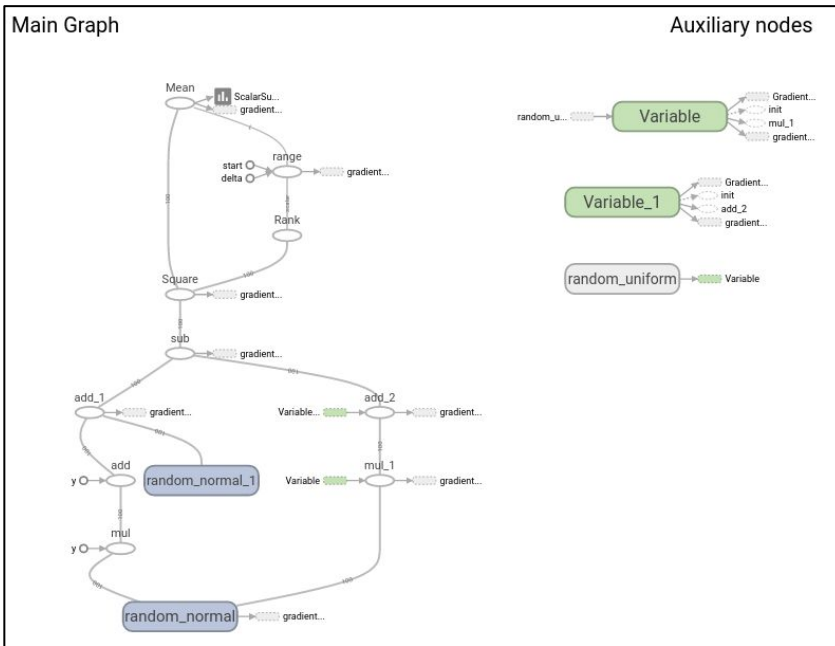
TensorBoard

- Some useful functions
 - `tf.name_scope('name')` → used to define a hierarchy of the nodes in the graph
 - Summary Operations ([Python API](#))
 - `tf.scalar_summary('name', values)` → summary for scalar values
 - `tf.histogram_summary('name', values)` → summary with a histogram of the values
 - `tf.merge_all_summaries()` → merges all summaries collected in the graph
 - Adding Summaries to Event Files ([Python API](#))
 - `tf.train.SummaryWriter('logdir', sess.graph)` → creates event file in the given directory
 - `.add_summary(summary, step)` → adds summary to the event file
- Most functions have a *name* parameter, use it for better graph organization
- You can organize different runs and compare them modifying *'logdir'*

TensorBoard

- Step 1: Name variables and make use of scopes for organizing your graph
- Step 2: Place summaries for those values you want to keep track
- Step 3: Let TF manage all summaries with `tf.merge_all_summaries()`
- Step 4: Create a writer pointing to the desired directory
- Step 5: Add the merged summaries to the writer
- Step 6: Launch TensorBoard

TensorBoard



```
import tensorflow as tf
```

```
tf.set_random_seed(1234)
```

```
x = tf.random_normal([100], mean=0.0, stddev=0.9)
```

```
y = x * 0.1 + 0.3 + tf.random_normal([100], mean=0.0, stddev=0.05)
```

```
W = tf.Variable(tf.random_uniform([], minval=-1.0, maxval=1.0))
```

```
b = tf.Variable(tf.zeros([]))
```

```
y_hat = W * x + b
```

```
loss = tf.reduce_mean(tf.square(y - y_hat))
```

```
train = tf.train.GradientDescentOptimizer(0.05).minimize(loss)
```

```
init = tf.initialize_all_variables()
```

```
sess = tf.Session()
```

```
writer = tf.train.SummaryWriter('/tmp/regression/', sess.graph)
```

```
sess.run(init)
```

```
for step in range(1, 101):
```

```
_, slope, intercept, error = sess.run([train, W, b, loss])
```

```
if step % 10 == 0:
```

```
    print('Step %.3d ; W = %.5f ; b = %.5f; loss = %.5f' %
          (step, slope, intercept, error))
```

TensorBoard

- Step 1: Name variables and make use of scopes for organizing your graph

```
with tf.name_scope('data'):
    with tf.name_scope('x'):
        x = tf.random_normal([100], mean=0.0, stddev=0.9, name='rand_x')
    with tf.name_scope('y'):
        y_true = x * tf.constant(0.1, name='real_slope') + tf.constant(0.3, name='bias') +
            tf.random_normal([100], mean=0.0, stddev=0.05, name='rand_y')

with tf.name_scope('W'):
    W = tf.Variable(tf.random_uniform([], minval=-1.0, maxval=1.0))

with tf.name_scope('b'):
    b = tf.Variable(tf.zeros([]))

with tf.name_scope('function'):
    y_pred = W * x + b

with tf.name_scope('error'):
    loss = tf.reduce_mean(tf.square(y_pred - y_true))
```

TensorBoard

- Step 2: Place summaries for those values you want to keep track

```
with tf.name_scope('W'):
    W = tf.Variable(tf.random_uniform([], minval=-1.0, maxval=1.0))
    tf.scalar_summary('function/W', W)

with tf.name_scope('b'):
    b = tf.Variable(tf.zeros([]))
    tf.scalar_summary('function/b', b)

with tf.name_scope('error'):
    loss = tf.reduce_mean(tf.square(y_pred - y_true))
    tf.scalar_summary('error', loss)
```

TensorBoard

- Step 3: Let TF manage all summaries with `tf.merge_all_summaries()`

```
merged = tf.merge_all_summaries()
```

- Step 4: Create a writer pointing to the desired directory
 - Use different paths for each run to compare the performance when modifying hyperparameters

```
writer = tf.train.SummaryWriter('/tmp/regression/run1', sess.graph)
```

- Step 5: Add the merged summaries to the writer
 - Specify a global step (iteration, etc.) to be added to the summary

```
summary_str = sess.run(merged)  
writer.add_summary(summary_str, step)
```

TensorBoard


- Step 6: Launch TensorBoard

Use the command → `$ tensorboard --logdir=/tmp/regression/`


Then browse → <http://localhost:6006/>

- If using Docker you may need to add the parameter `-p 6006:6006`

Main Graph




Auxiliary nodes



TenosorBoard

- Some runtime statistics in TensorBoard
 - See and run `regression_tb_md.py` for basic changes
 - Also available in [Running demo \(code\)](#)
- More advanced CPU and GPU tracing
 - Steps to follow to enable it → [Link](#)
 - Also, lines commented in `regression_tb_md.py`
 - Go to `chrome://tracing` and load the json file obtained



Example of a tracing with one GPU (not the regression code)

TensorBoard

- Homework:
 - Build a TB version of the simple network build today
 - Make use of `tf.histogram_summary('name', values)` for non-scalar variables